



# Bahria University

Department of Computer Science

## Natural Language Processing

CSC 441

**Project**

**Final Deliverable**

---

**Hannan Ahmed Jadran 01-134211-106**

**Muhammad Ali Momin 01-134211-051**

**BS-CS-7B**

---

**Submitted to**

**Dr. Arifur Rahman**

**Head of Department / Professor**

# Contents

Abstract.....	3
1. Introduction .....	3
2. Methodology .....	4
2.1. Preprocessing .....	4
Preprocessing steps .....	5
2.1. Feature Extraction.....	6
TF-IDF (Term Frequency-Inverse Document Frequency) Vectorization .....	6
2.2. Application development.....	6
2.3. Data Ingestion .....	7
2.4. Data Cleaning .....	7
2.5. Named Entity Recognition .....	7
2.6. Post Processing.....	8
2.7. Document Generation.....	8
2.8. Error Handling and Validation .....	8
3. Dependencies.....	9
• Flask.....	9
• NLTK (Natural Language Toolkit) .....	9
• re (Regular Expressions) .....	9
• python-docx.....	9
4. Application and Deployment.....	9
Links:.....	9
References .....	9

# Figures

Figure 1 General Flow of Operations.....	3
Figure 2 Methodology steps .....	4
Figure 3: Preprocessing Steps .....	5
Figure 4: Feature Extraction.....	6
Figure 5: Application Workflow .....	6
Figure 6: Data Ingestion System Diagram .....	7
Figure 7: Data Cleaning Flowchart .....	7
Figure 8: NER Module Processing .....	7
Figure 9: Document Generation Module .....	8
Figure 10: Error Handling Flowchart .....	8

# Multilingual Named Entity Extraction System

## Abstract

Named Entity Recognition (NER) is a crucial aspect of Natural Language Processing (NLP) that involves identifying and classifying key information (entities) within a body of text. This report outlines a system designed to extract named entities from PDF documents, translate them if necessary, and save the processed text in a Word document with highlighted entities. This system leverages SpaCy for NER, PyMuPDF for PDF handling, and Google Translate for language translation.

## 1. Introduction

Named entity recognition (NER) is a fundamental task in natural language processing (NLP) that involves identifying and classifying named entities, such as persons, organizations, locations, and other relevant entities, within unstructured text data. NER has numerous applications, including information extraction, question answering, and knowledge base construction.

In multilingual environments, NER becomes more challenging due to the need to handle texts in different languages. This project focuses on developing an NLP system that can extract named entities from documents written in Urdu or English. Urdu, being a widely spoken language in South Asia, presents unique challenges due to its complex script and linguistic properties.

The proposed system employs a two-step approach for Urdu documents. First, the Urdu text is translated into English using machine translation techniques. This step is necessary because most state-of-the-art NER models are primarily trained on English data. Once the text is in English, a pretrained NER model is applied to identify, and extract named entities from the translated text. After the named entities have been extracted from the English text, they are converted back into their original Urdu script. This step ensures that the output is consistent with the input document's language, providing a seamless user experience for Urdu speakers. For English documents, the system directly applies the NER model to the input text, bypassing the translation step.

This project aims to provide a robust and efficient solution for named entity extraction in both Urdu and English documents, enabling applications such as information retrieval, text mining, and knowledge management in multilingual environment.

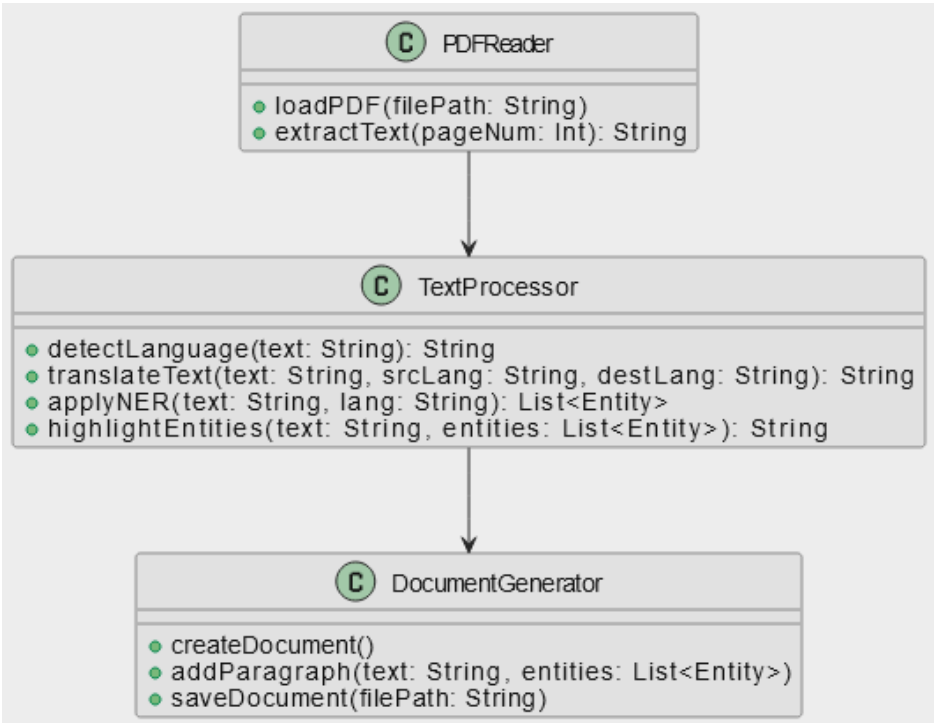


Figure 1 General Flow of Operations

## 2. Methodology

The methodology section provides an in-depth explanation of each step involved in developing and deploying the NER application. This includes the design of the system architecture, data processing, and the implementation of key functionalities.

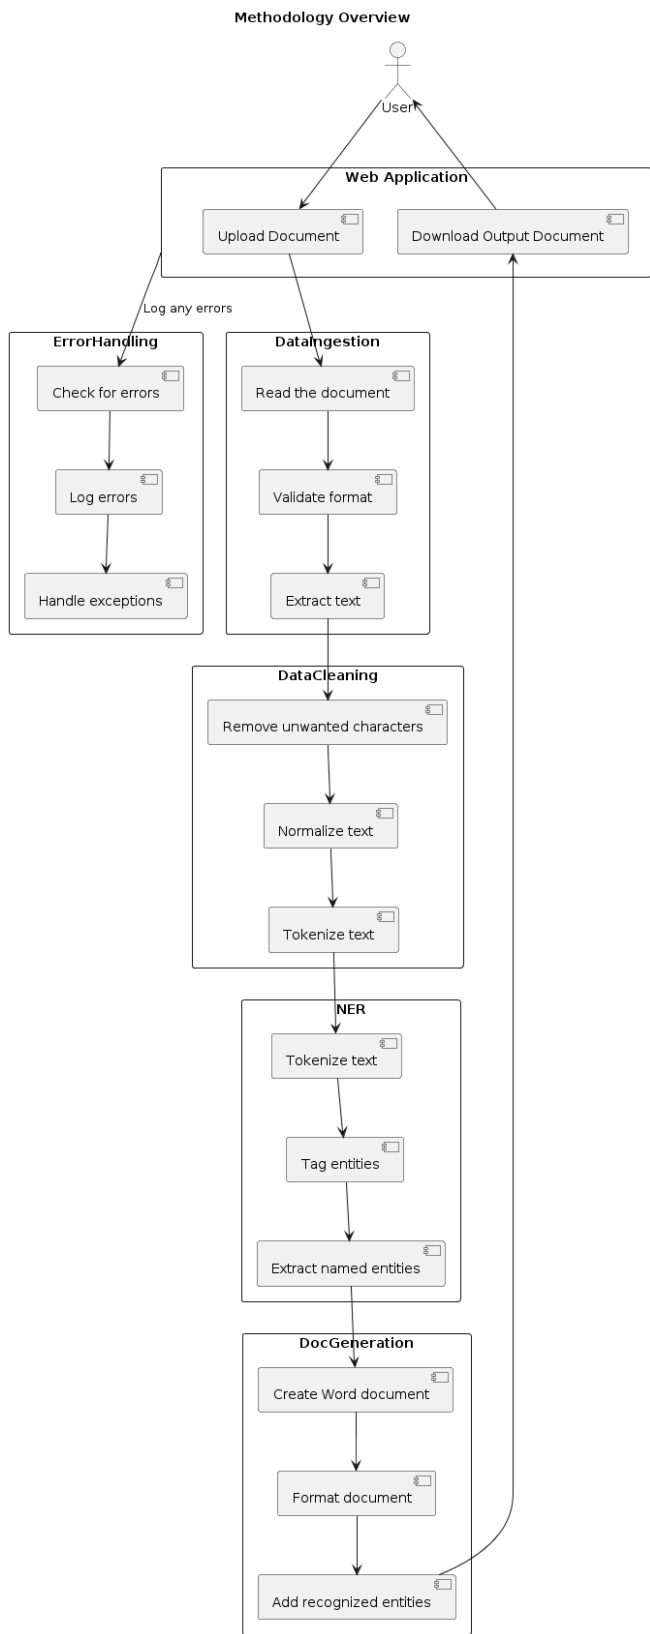
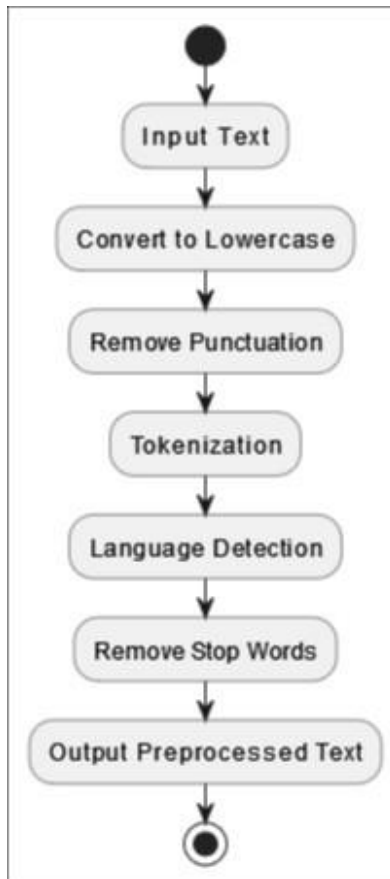


Figure 2 Methodology steps

### 2.1. Preprocessing

The preprocessing stage serves as a pivotal preparatory step, laying the foundation for subsequent text analysis tasks. Its primary goal is to refine raw text data, ensuring its compatibility with downstream processes like named entity recognition (NER) and other natural language processing (NLP) tasks.



*Figure 3: Preprocessing Steps*

## Preprocessing steps

### I. Convert to Plain Text

Initially, the text undergoes a transformation to strip it of any extraneous elements, ensuring it exists in a clean, plain text format devoid of formatting artifacts, such as leading or trailing whitespace, or any other non-textual elements that could obfuscate analysis.

### II. Check Language

Following the conversion to plain text, an important consideration is determining the language of the text. This step is pivotal, especially for multilingual applications, as it allows for the application of language-specific preprocessing techniques tailored to the linguistic nuances of the text. For instance, distinct preprocessing strategies might be employed for Urdu and English text due to differences in syntax, vocabulary, and structure.

### III. Preprocessing and Clean Data

- Once the language is identified, the text is subjected to a series of cleaning operations aimed at refining its quality and enhancing its suitability for analysis. This multifaceted process encompasses several key operations:
- Lowercasing: The text is uniformly converted to lowercase to facilitate standardization and mitigate issues arising from case discrepancies during subsequent analysis.
- Punctuation and Special Character Removal: Extraneous elements such as punctuation marks and special characters are systematically stripped from the text. This serves to declutter the data and eliminate non-semantic elements that could potentially interfere with analysis.
- Tokenization: The text is segmented into individual tokens or words, a foundational step for subsequent analysis tasks. Tokenization serves to break down the text into its constituent units, enabling granular analysis at the word level.
- Stop Word Removal: Stop words, commonly occurring words that contribute little to the overall semantic content of the text, are pruned from the data. This helps streamline analysis by focusing on substantive content while discarding superfluous linguistic elements.

## 2.1. Feature Extraction

The objective of feature extraction is to transform the preprocessed text data into a structured format that can be utilized as input for subsequent analysis tasks. By extracting relevant features, we aim to capture the essential information encoded within the text, enabling effective analysis and interpretation.

### TF-IDF (Term Frequency-Inverse Document Frequency) Vectorization

TF-IDF is a widely used technique for feature extraction in text analysis. It calculates the importance of a term within a document relative to a corpus of documents. The TF-IDF value increases proportionally to the number of times a term appears in the document but is offset by the frequency of the term in the corpus, which helps to adjust for the fact that some terms are more common overall.

- Implementation of feature extraction techniques typically involves utilizing libraries or tools that provide implementations of these techniques. For example:
- Scikit-learn provides a TfidfVectorizer class for TF-IDF vectorization.
- Domain-specific feature engineering approaches may require custom implementation based on the specific requirements of the analysis task.
- The preprocessed text data serves as input to these feature extraction techniques, and the resulting feature representations are used for subsequent analysis tasks such as classification, clustering, or information retrieval.

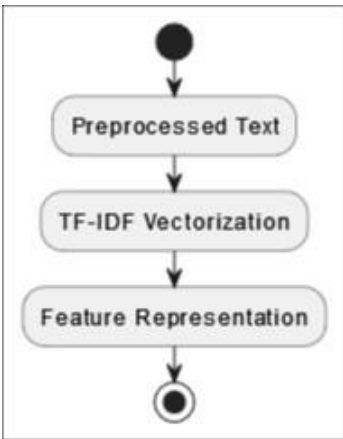


Figure 4: Feature Extraction

## 2.2. Application development

The system is built using a modular approach, where each module is responsible for a specific task. The main components include the Flask web application, the NER module, and the document processing module. The Flask application serves as the interface between the user and the backend processing logic.




Figure 5: Application Workflow

### 2.3. Data Ingestion

The first step in the process is data ingestion. The user uploads a text file through the web interface. The Flask application handles this file upload, ensuring that the file is correctly received and stored temporarily for processing.




Figure 6: Data Ingestion System Diagram

### 2.4. Data Cleaning

Once the file is uploaded, the text content is extracted and passed to the data cleaning module. This step involves removing unwanted characters, such as control characters or non-ASCII characters, that may interfere with the processing. The text is then standardized to ensure consistency in the subsequent analysis.

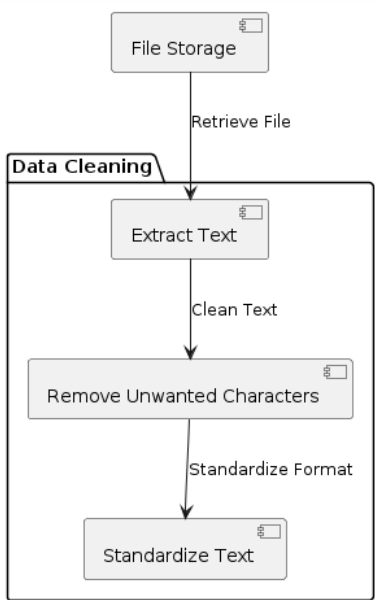


Figure 7: Data Cleaning Flowchart

### 2.5. Named Entity Recognition

The core of the application is the NER module. This module uses pre-trained machine learning models to identify and classify named entities within the text. The models are trained on large corpora of labeled data, allowing them to recognize patterns and classify entities with high accuracy. The NER process involves tokenizing the text, applying the model, and extracting the entities along with their corresponding types.

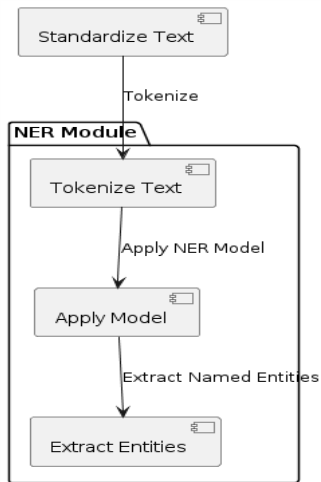


Figure 8: NER Module Processing

## 2.6. Post Processing

After extracting the named entities, the text undergoes post-processing to format the output appropriately. This includes reassembling the text with highlighted or annotated named entities, which makes it easier for the user to identify the relevant information.

## 2.7. Document Generation

The final processed text is then used to generate a Word document. The python-docx library is utilized for this purpose. The module creates a new document, adds the processed text, and ensures that the formatting is preserved. The document is saved and made available for the user to download.

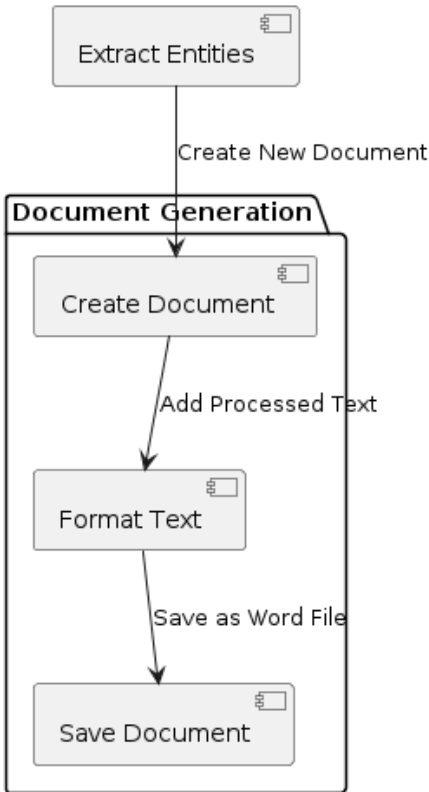


Figure 9: Document Generation Module

## 2.8. Error Handling and Validation

Throughout the process, the application incorporates robust error handling and validation mechanisms. This ensures that any issues encountered during file upload, text processing, or document generation are properly managed and communicated to the user.

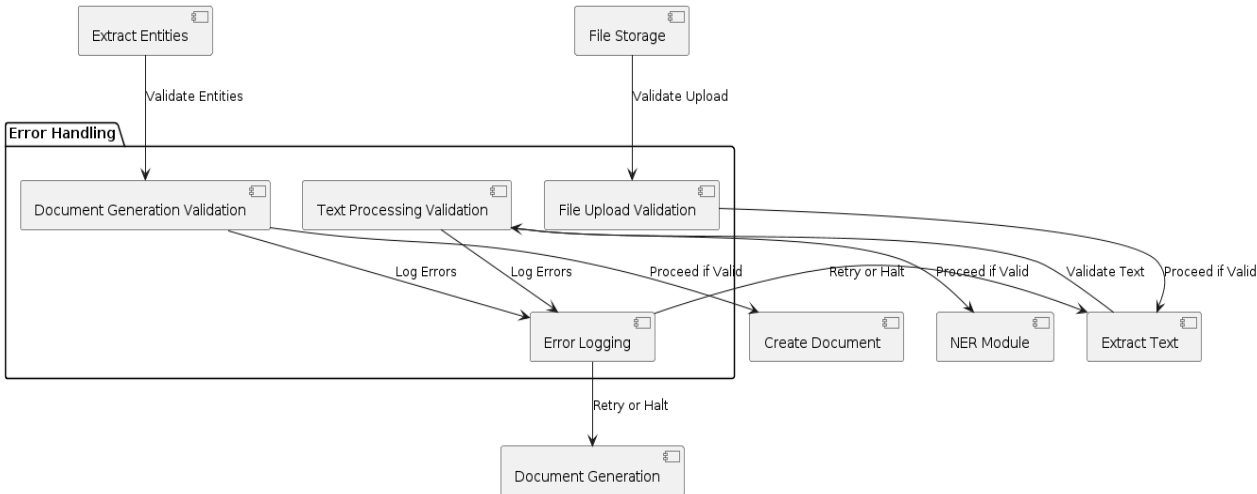


Figure 10: Error Handling Flowchart



### 3. Dependencies

- **Flask**

Flask is a lightweight WSGI web application framework in Python. It is designed to make getting started quick and easy, with the ability to scale up to complex applications.

- **NLTK (Natural Language Toolkit)**

NLTK is a leading platform for building Python programs to work with human language data. It provides easy-to-use interfaces to over 50 corpora and lexical resources, such as WordNet, along with a suite of text processing libraries.

- **re (Regular Expressions)**

The re module provides regular expression matching operations similar to those found in Perl. It is used for text searching and manipulation.

- **python-docx**

python-docx is a Python library for creating and updating Microsoft Word (.docx) files. It allows for complex document generation and formatting.

### 4. Application and Deployment

The application is deployed on a web server, allowing users to access it via a web browser. The Flask application is configured to run on the server, handling incoming requests and serving the user interface. The deployment process involves setting up the server environment, installing necessary dependencies, and configuring the web server to forward requests to the Flask application.

The deployment ensures that the application is accessible to users, with a focus on scalability and reliability. The server setup includes load balancing and failover mechanisms to handle multiple user requests simultaneously and ensure continuous availability.

### Links:

**Linkedin:** <https://www.linkedin.com/feed/update/urn:li:activity:7198031288887324673/>

**Github:** <https://github.com/gibran404/NLP>

### References

Beautiful Soup Documentation - <https://www.crummy.com/software/BeautifulSoup/bs4/doc/>

Scikit-learn Documentation - <https://scikit-learn.org/stable/>

Eftekhari, A. (2020). Machine learning for named entity recognition: A study on dataset selection and model confidence.

Mahmood, A., Hussain, F., Shafiq, H. M., C Siddiqui, E. (2010, October). Machine translation between English and Urdu. In 2010 International Conference on Arabic and Multilingual Information Processing (pp. 105-114). IEEE.

Leaman, R., Islamaj Doğan, R., C Lu, Z. (2020). DNorm: disease name normalization with pairwise vector space mappings. Journal of the American Medical Informatics Association, 27(Supplement\_1), 49-55.

Lample, G., Ballesteros, M., Subramanian, S., Kawakami, K., C Dyer, C. (2016). Neural architectures for named entity recognition.

Riza Batista-Navarro, R. T., Rak, R., C Ananiadou, S. (2015, July). Analysis of live multi- domain machine translation for multilingual named entity recognition. In Proceedings of the 11th International Conference on Computational Semantics (pp. 68-76).